

UNIVERSITY OF ST ANDREWS

COMPUTER GRAPHICS


CS4102

Ring-Based Distributed System

Author:

150008022

May 15, 2020



Goal

The aim of this practical is to understand the key principles behind various techniques frequently used for the rendering of 3D objects, and to get hands-on experience with their implementation and manipulation.

1 Running the Program

The project uses `maven` for lifecycle and dependency management. `LWJGL` was used for this project, which provides a way to use `OpenGL` from Java. `OpenGL` *may* be a cause of compatibility issues when running, but an early and forward compatible profile was used to minimise the risk of this.

Usage

`mvn clean install` will compile the project, run all tests, and produce an executable jar.

`cd target` to enter the directory where the jar has been compiled.

`java -jar FaceModelling-jar-with-dependencies.jar` to run the program. Pass the `-h` flag to see options.

2 Implementation

2.1 Loading Mesh Data

A `FaceLoader` instance initialises by parsing the indices (`mesh.csv`), and the average face coordinates (`sh_000.csv`), vertex weighting (`sh_ev.csv`), average face colour (`tx_000.csv`), and colour weightings (`tx_ev.csv`). The indices are converted from 1-indexed to 0-indexed when being loaded in. The `FaceLoader` instance then loads the deltas for both coordinates and colours for each face, calculating the actual coordinates for each vertex of each face.

The arrays of face data are then passed to the `MeshLoader` which creates a vertex array object (VAO) for each face, and binds a vertex buffer object (VBO) for the vertices, indices, and colours to them.

The use of indices instead of simply listing all coordinates saves on memory as the one vertex can be used in many triangles, as demonstrated in figure 1.




Figure 1: How indices are used when describing meshes.

2.2 Depth Testing and Painters Algorithm

OpenGL does not natively implement painters algorithm for depth testing, and does not expose an API for carrying it out. Instead, it uses Z-buffers, Z-Buffers involve recording the z-index last drawn at each pixel of the window view. A pixel with a greater z-index than the current value will be drawn and the z-buffer updated to reflect the new value. Otherwise the depth test fails and the object will not be drawn at that coordinate. This functionality is enabled using the `glEnable(GL11.GL_DEPTH_TEST)` call. An optimisation called early z-testing is often used which is applied earlier on in the rendering pipeline in order to reduce the number of fragments that need to be processed in total.

Painters algorithm and Z-buffers are both viable solutions to the same problem; the former is more computationally expensive while the latter requires more memory. Z-Buffers are preferred in most cases for modern hardware since memory has become cheaper. Painters algorithm also requires additional computation in the case where objects intersect. Methods to solve this involve splitting the intersecting shapes and rendering these instead.

A separate renderer was implemented that used painters algorithm. This

renderer was significantly slower both due to the added computation necessary for the painters algorithm, and also due to the fact that the sorting was performed using the CPU rather than the GPU.

The first step involved applying each transformation to the set of coordinates. To try improve performance, this task was parallelised using a thread pool. Then, for each triplet of indices (i.e. each triangle), the maximum Z value was recorded. Next, maximum z values for each triangle were sorted lowest to highest using a custom `quicksort` implementation. The custom implementation was necessary so that every swap performed on the array of maximum z values would also be performed on the indices array, effectively sorting each triangle by maximum z value. The new indices would then be passed as a VBO for OpenGL, which renders triangles in order of appearance in the indices VBO.

2.3 Selection Triangle

The selection triangle is rendered with the three control faces at each vertex. A left click on the triangle updates the weighting to be used when producing the interpolated face. This required carrying out ray casting with respect to the mouse position, and was achieved using the steps described by [1]. Mapping the left click to a position on the selection triangle involved converting from the 2D screen space coordinates to OpenGLs 3D device coordinate system (Figure 2). We then apply the inverse of the transformation matrices we used to convert to screen space to get the coordinates in world space.

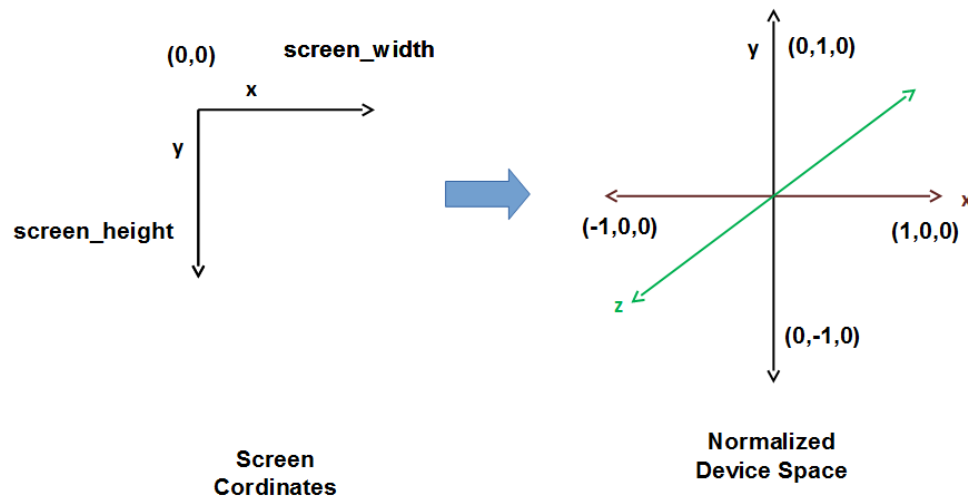


Figure 2: Converting from screen space to OpenGL [2].

3 Summary of Implementation

Requirements

- Render faces using index, coordinate, and colour buffers
- Painters algorithm
- Click triangle to select weighting between different faces
- Interpolate to create face from control faces with weighting
- Flat shading
- Gouraud shading
- Orthographic projection
- Perspective projection
- Adjustable focal length
- Comparison of shading and projection
- Light attention model from lectures
- Adjustable free parameters
- Phongs secularity term in reflectance model

Extensions

- Camera is able to move around 3D world
- Comparison between painters algorithm and z-buffers
- Control faces are also visible during selection
- Use of OpenGL for optimised rendering
- Scene graph to render children objects relative to parents

References

- [1] Anton Gerdelan. *Mouse Picking with Ray Casting*. <https://antongerdelan.net/opengl/raycasting.html>.
- [2] Antonio Hernández Bejarano. *3D Game Development with LWJGL*. <https://lwjglgamedev.gitbooks.io/3d-game-development-with-lwjgl>.